

Telecom Customer Churn Prediction

Phase 3 Machine Learning Project — CRISP-DM Framework

Author: Data Science Student

Dataset: Telecom Customer Churn Dataset (3,333 records, 21 features)

Target Variable: churn (Boolean — True if customer cancelled service)

Project methodology

1. Business Understanding

The objective of this project is to predict customer churn in a telecom company. By identifying customers likely to leave, the business can take proactive retention measures, reduce revenue loss, and improve customer satisfaction.

Key Goal: Build a predictive model to identify churn and derive actionable insights.

2. Data Understanding

In this section, we explore the dataset to understand its structure, variables, and quality.

Focus Areas:

- Data types and distributions
- Missing values
- Initial patterns and anomalies

3. Data Preparation

This stage involves cleaning and transforming data for modeling.

Steps include:

- Handling missing values
- Encoding categorical variables
- Feature scaling
- Feature selection

Important: Clearly justify why features are dropped or transformed.

4. Data Analysis

EDA is used to uncover patterns, relationships, and insights in the data.

For every visualization, ensure interpretation includes:

- What the pattern shows
- Why it matters
- How it influences modeling decisions

Examples of insights to highlight:

- Class imbalance in churn
- Strong predictors (e.g., tenure, charges)
- Correlations between variables

5. Modeling

Multiple models are trained to compare performance.

Models may include:

- Logistic Regression
- Decision Tree
- Random Forest

Enhancements:

- Hyperparameter tuning (e.g., GridSearchCV)
- Cross-validation

6. Evaluation

Models are evaluated using appropriate metrics such as:

- Accuracy
- Precision
- Recall
- F1-score

Key Requirement:

- Compare models clearly
- Explain why the best model performs better
- Relate results back to business impact (churn reduction)

7. Conclusion & Recommendations

Summarize key findings and provide business recommendations.

Include:

- Key drivers of churn
- Best performing model
- Actionable strategies (e.g., targeting high-risk customers)

Business Value: Explain how the model can help reduce churn and improve profitability.

Phase 1: Business Understanding

Real-World Problem

Customer churn — when a customer stops doing business with a company — is one of the most expensive problems in telecommunications. Acquiring a new customer costs **five to seven times more** than retaining an existing one. For a telecom provider managing thousands of accounts, even a modest improvement in retention has a direct, material impact on revenue.

Stakeholder

The **Customer Retention & Marketing team** at a mid-sized US telecom operator. This team is responsible for proactively reaching out to at-risk customers with targeted offers, service upgrades, or loyalty incentives *before* those customers cancel.

Why Machine Learning?

A simple rule such as "flag customers with more than 3 service calls" captures one signal but misses the combined, non-linear interactions between usage patterns, plan type, call volume, and account history that a classification model can capture and quantify.

Business Success Metric

The cost of a **false negative** (missing a real churner) is high — the customer is lost. The cost of a **false positive** (contacting a loyal customer) is low — a wasted retention offer.

Primary metric: Recall for the churned class — catch as many true churners as possible.

Secondary metric: F1-score — balance recall against precision.

Tertiary metric: ROC-AUC — evaluate overall ability to rank customers by churn risk.

Phase 2: Data Understanding

```
In [3]: # Importing Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.patches as mpatches
import matplotlib.gridspec as gridspec
import seaborn as sns
import warnings
!pip install squarify
import squarify

warnings.filterwarnings('ignore')
sns.set_style('whitegrid')
plt.rcParams.update({
    'figure.dpi': 120,
    'axes.titlesize': 13,
    'axes.titleweight': 'bold',
    'axes.spines.top': False,
    'axes.spines.right': False
})

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import (
    classification_report, confusion_matrix, ConfusionMatrixDisplay,
    roc_auc_score, recall_score, precision_score, f1_score,
    roc_curve, precision_recall_curve, average_precision_score
)

RANDOM_STATE = 42
# Consistent colour palette used throughout
C_RETAIN = '#2ecc71' # green - retained customers
C_CHURN = '#e74c3c' # red - churned customers
C_BLUE = '#3498db' # blue - train / model 1
C_ORANGE = '#e67e22' # orange - model 2
C_PURPLE = '#9b59b6' # purple - model 5 / final

print('All libraries imported successfully.')
```

Collecting squarify
 Downloading squarify-0.4.4-py3-none-any.whl (4.1 kB)
 Installing collected packages: squarify
 Successfully installed squarify-0.4.4
 All libraries imported successfully.

```
In [4]: # Loading the dataset and identifying the number of rows and columns
df = pd.read_csv('churn_telecoms_dataset.csv')
print(f'Dataset shape: {df.shape[0]:,} rows x {df.shape[1]} columns')
df.head()
```

Dataset shape: 3,333 rows x 21 columns

Out[4]:

	state	account length	area code	phone number	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total day charge	...	total eve calls	tc
0	KS	128	415	382-4657	no	yes	25	265.1	110	45.07	...	99	16
1	OH	107	415	371-7191	no	yes	26	161.6	123	27.47	...	103	16
2	NJ	137	415	358-1921	no	no	0	243.4	114	41.38	...	110	10
3	OH	84	408	375-9999	yes	no	0	299.4	71	50.90	...	88	5
4	OK	75	415	330-6626	yes	no	0	166.7	113	28.34	...	122	12

5 rows × 21 columns



In [5]:

```
# Determining the data types and missing values of the dataset
print('=== Data Types ===')
print(df.dtypes.to_string())
print(f'\n=== Missing Values (Total: {df.isnull().sum().sum()}) ===')
print(df.isnull().sum().to_string())
```

```
=== Data Types ===
state                object
account length      int64
area code           int64
phone number        object
international plan   object
voice mail plan      object
number vmail messages  int64
total day minutes    float64
total day calls      int64
total day charge     float64
total eve minutes    float64
total eve calls      int64
total eve charge     float64
total night minutes  float64
total night calls    int64
total night charge   float64
total intl minutes   float64
total intl calls     int64
total intl charge    float64
customer service calls int64
churn                bool
```

```
=== Missing Values (Total: 0) ===
state                0
account length      0
area code           0
phone number        0
international plan   0
voice mail plan      0
number vmail messages 0
total day minutes    0
total day calls      0
total day charge     0
```

```

total eve minutes      0
total eve calls        0
total eve charge       0
total night minutes    0
total night calls      0
total night charge     0
total intl minutes     0
total intl calls       0
total intl charge      0
customer service calls 0
churn                  0

```

In [6]: `df.describe().round(2)`

Out[6]:

	account length	area code	number vmail messages	total day minutes	total day calls	total day charge	total eve minutes	total eve calls	total eve charge	total night minutes	total night calls
count	3333.00	3333.00	3333.00	3333.00	3333.00	3333.00	3333.00	3333.00	3333.00	3333.00	3333.00
mean	101.06	437.18	8.10	179.78	100.44	30.56	200.98	100.11	17.08	200.87	100.44
std	39.82	42.37	13.69	54.47	20.07	9.26	50.71	19.92	4.31	50.57	19.92
min	1.00	408.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	23.20	33.00
25%	74.00	408.00	0.00	143.70	87.00	24.43	166.60	87.00	14.16	167.00	87.00
50%	101.00	415.00	0.00	179.40	101.00	30.50	201.40	100.00	17.12	201.20	100.44
75%	127.00	510.00	20.00	216.40	114.00	36.79	235.30	114.00	20.00	235.30	113.00
max	243.00	510.00	51.00	350.80	165.00	59.64	363.70	170.00	30.91	395.00	175.00

2.1 Target Variable — Waffle Chart

A waffle chart gives a more intuitive feel for the class imbalance than a pie chart: each square represents 1% of customers.

```

In [7]: # VISUAL 1: Waffle chart - churn rate
churn_rate = df['churn'].mean() * 100 # 14.49 %
churned_sq = round(churn_rate) # 14 red squares

fig, ax = plt.subplots(figsize=(12, 6))
ax.set_aspect('equal')

for i in range(100):
    row = i // 10
    col = i % 10
    color = C_CHURN if i < churned_sq else C_RETAIN
    rect = plt.Rectangle((col * 1.1, row * 1.1), 1.0, 1.0,
                        color=color, linewidth=1.5, edgecolor='white')
    ax.add_patch(rect)
    # Label each square with its percentage
    if i == 0 or i == churned_sq:
        pass

ax.set_xlim(-0.2, 11.2)
ax.set_ylim(-0.5, 11.5)

```

```

ax.axis('off')


patch_churn = mpatches.Patch(color=C_CHURN, label=f'Churned — {df["churn"].sum():,} c
patch_retain = mpatches.Patch(color=C_RETAIN, label=f'Retained — {(~df["churn"]).sum():,} c
ax.legend(handles=[patch_churn, patch_retain], loc='lower center',
          fontsize=13, frameon=False, ncol=2,
          bbox_to_anchor=(0.5, -0.05))

ax.set_title('Customer Churn Rate — Each Square = 1% of Customers',
            fontsize=15, fontweight='bold', pad=15)
plt.tight_layout()
plt.show()

print(f'Churn rate: {churn_rate:.2f}%')
print('A naive "predict no churn" model gets 85.5% accuracy but catches ZERO churners.')
print('→ Accuracy is a misleading metric here. We use Recall and F1.')

```

Customer Churn Rate — Each Square = 1% of Customers



■ Churned — 483 customers (14.5%)
 ■ Retained — 2,850 customers (85.5%)

Churn rate: 14.49%

A naive "predict no churn" model gets 85.5% accuracy but catches ZERO churners.

→ Accuracy is a misleading metric here. We use Recall and F1.

2.2 Usage Distributions — KDE Plots

Kernel Density Estimation plots show the full distribution shape of each usage feature, split by churn status — smoother and more informative than histograms.

```

In [8]: # VISUAL 2: KDE plots — usage by churn status
        usage_cols = [

```

```

('total day minutes', 'Total Day Minutes'),
('total eve minutes', 'Total Evening Minutes'),
('total night minutes', 'Total Night Minutes'),
('total intl minutes', 'Total International Minutes'),
('customer service calls', 'Customer Service Calls'),
('account length', 'Account Length (days)'),
]

fig, axes = plt.subplots(2, 3, figsize=(18, 10))
axes = axes.flatten()

for ax, (col, title) in zip(axes, usage_cols):
    for churned, color, label in [
        (False, C_RETAIN, f'Retained (mean={df[df["churn"]==False][col].mean():.1f})'),
        (True, C_CHURN, f'Churned (mean={df[df["churn"]==True][col].mean():.1f})')
    ]:
        subset = df[df['churn'] == churned][col]
        subset.plot.kde(ax=ax, color=color, linewidth=2.5, label=label)
        ax.fill_between(
            np.linspace(subset.min(), subset.max(), 300),
            0,
            [ax.lines[-1].get_ydata().max() * 0] * 300, # placeholder
            alpha=0
        )
        # Shade under each KDE
        line = ax.lines[-1]
        ax.fill_between(line.get_xdata(), line.get_ydata(), alpha=0.15, color=color)


    ax.set_title(title)
    ax.set_ylabel('Density')
    ax.set_xlabel('')
    ax.legend(fontsize=8)

plt.suptitle('Feature Distributions: Churned vs Retained Customers (KDE)',
            fontsize=16, fontweight='bold', y=1.01)
plt.tight_layout()
plt.show()

print('Key insight: Churned customers average 207 day minutes vs 175 retained.')
print('Higher daytime usage → higher bills → higher price sensitivity → more likely to

```


Feature Distributions: Churned vs Retained Customers (KDE)



Key insight: Churned customers average 207 day minutes vs 175 retained.
Higher daytime usage → higher bills → higher price sensitivity → more likely to leave.

2.3 Violin Plots — Spread and Shape by Churn Status

Violin plots combine a box plot with a KDE — showing median, IQR, and the full distribution shape in one graphic. Ideal for comparing how spread out usage patterns are between churned and retained customers.


```
In [9]: # VISUAL 3: Violin plots – charges by time-of-day and churn
df_plot = df[['total day minutes', 'total eve minutes',
              'total night minutes', 'total intl minutes', 'churn']].copy()
df_melt = df_plot.melt(id_vars='churn', var_name='Usage Type', value_name='Minutes')
df_melt['Churn Status'] = df_melt['churn'].map({False: 'Retained', True: 'Churned'})
df_melt['Usage Type'] = df_melt['Usage Type'].str.replace('total ', '').str.replace(' m', '')

fig, ax = plt.subplots(figsize=(15, 7))

sns.violinplot(
    data=df_melt,
    x='Usage Type', y='Minutes', hue='Churn Status',
    palette={'Retained': C_RETAIN, 'Churned': C_CHURN},
    split=True, inner='quartile', linewidth=1.2,
    ax=ax
)

ax.set_title('Usage Minutes Distribution by Time of Day and Churn Status\n'
            '(Split violin: left = Retained, right = Churned | Inner lines = quartiles)'
            , fontsize=13, fontweight='bold')
ax.set_xlabel('Usage Period', fontsize=12)
ax.set_ylabel('Minutes', fontsize=12)
ax.legend(title='Churn Status', fontsize=11)
plt.tight_layout()
plt.show()
```

```
print('Day minutes show the largest separation between churned and retained customers.')
print('Night and evening minutes show very similar distributions – less predictive.')
```



Day minutes show the largest separation between churned and retained customers.
Night and evening minutes show very similar distributions – less predictive.

2.4 Customer Service Calls — Strip Plot with Churn Rate Overlay

A strip plot shows individual data points, revealing the raw data density at each call count level. Overlaying the churn rate line connects the individual observations to the aggregate pattern.

```
In [10]: # VISUAL 4: Strip plot + churn rate line – customer service calls
fig, (ax1, ax2) = plt.subplots(2, 1, figsize=(13, 10), sharex=True,
                               gridspec_kw={'height_ratios': [2, 1]})

# Top: strip plot of individual customers
df_strip = df[['customer service calls', 'churn']].copy()
df_strip['Churn Status'] = df_strip['churn'].map({False: 'Retained', True: 'Churned'})

sns.stripplot(
    data=df_strip, x='customer service calls', y='Churn Status',
    palette={'Retained': C_RETAIN, 'Churned': C_CHURN},
    alpha=0.25, jitter=0.3, size=4, ax=ax1
)
ax1.set_title('Individual Customers by Service Calls (top) and Churn Rate (bottom)',
              fontsize=13, fontweight='bold')
ax1.set_xlabel('')
ax1.set_ylabel('Churn Status', fontsize=11)

# Bottom: churn rate line
csc = df.groupby('customer service calls')['churn'].agg(['mean', 'count']).reset_index()
csc['mean'] = csc['mean'] * 100

ax2.plot(csc['customer service calls'], csc['mean'],
         color=C_CHURN, linewidth=2.5, marker='o', markersize=8, zorder=5)
ax2.fill_between(csc['customer service calls'], csc['mean'],
                alpha=0.15, color=C_CHURN)
ax2.axhline(y=churn_rate, color='navy', linestyle='--', linewidth=1.5,
            label=f'Overall churn rate: {churn_rate:.1f}%')
ax2.axvspan(3.5, 9.5, alpha=0.08, color=C_CHURN, label='High-risk zone (4+ calls)')
```

```

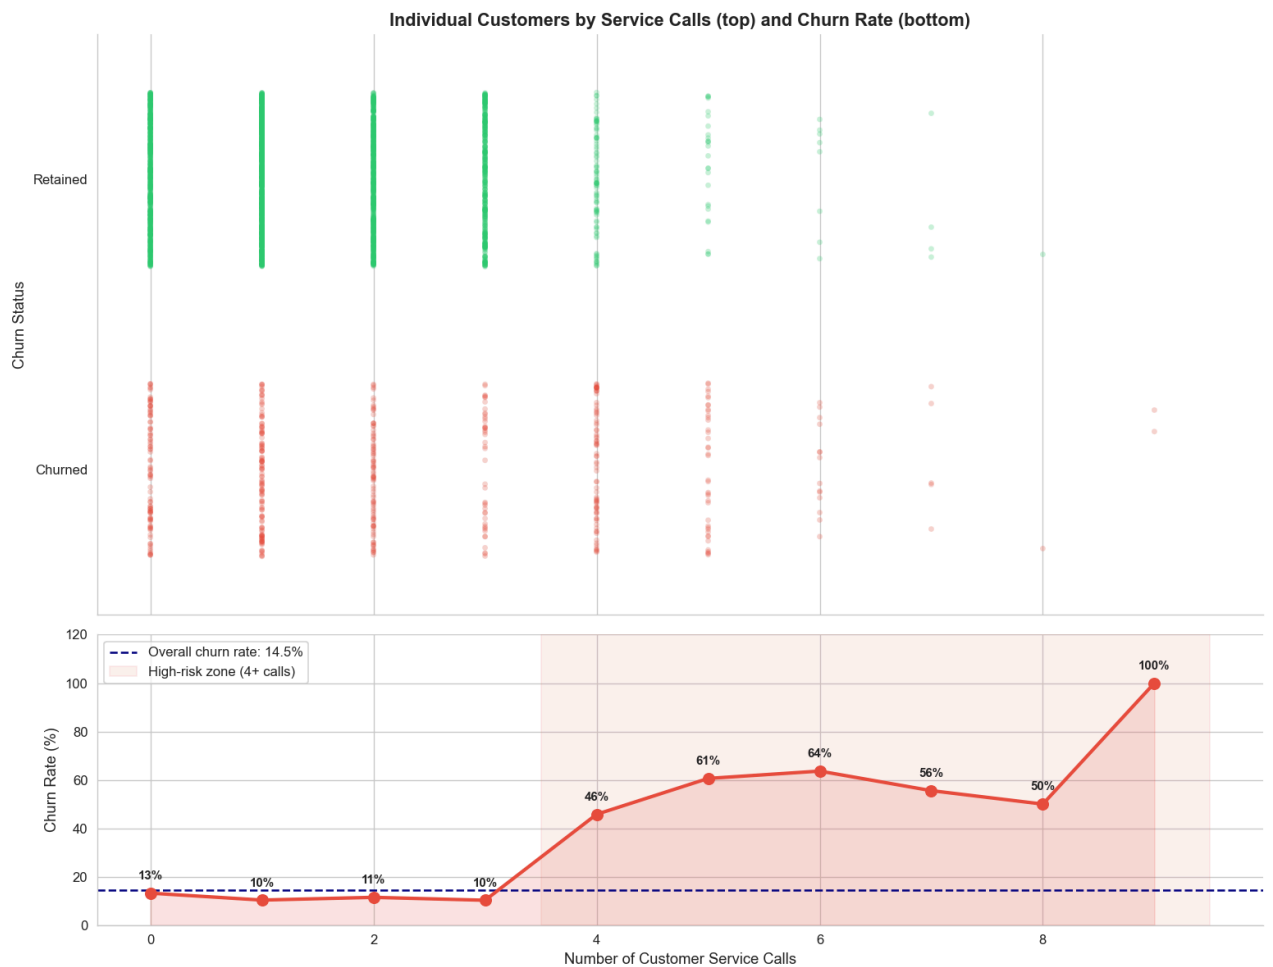
for _, row in csc.iterrows():
    ax2.annotate(f'{row["mean"]:.0f}%',
                (row['customer service calls'], row['mean']),
                textcoords='offset points', xytext=(0, 10),
                ha='center', fontsize=9, fontweight='bold')

ax2.set_xlabel('Number of Customer Service Calls', fontsize=11)
ax2.set_ylabel('Churn Rate (%)', fontsize=11)
ax2.legend(fontsize=10)
ax2.set_ylim(0, 120)

plt.tight_layout()
plt.show()

low = df[df['customer service calls'] <= 3]['churn'].mean() * 100
high = df[df['customer service calls'] >= 4]['churn'].mean() * 100
print(f'Churn rate 0-3 service calls: {low:.1f}%')
print(f'Churn rate 4+ service calls: {high:.1f}% (~{high/low:.1f}x higher)')

```



Churn rate 0-3 service calls: 11.3%

Churn rate 4+ service calls: 51.7% ($\approx 4.6\times$ higher)

2.5 Pair Plot — Top Feature Interactions

A pair plot reveals pairwise relationships between the top predictive features, coloured by churn status.

```

In [12]: # VISUAL 5: Pair plot - top features coloured by churn
top_features = ['total day minutes', 'customer service calls',

```


```

'total intl minutes', 'number vmail messages']

df_pair = df[top_features + ['churn']].copy()
df_pair['Churn'] = df_pair['churn'].map({False: 'Retained', True: 'Churned'})

pair_grid = sns.pairplot(
    df_pair.drop(columns='churn'),
    hue='Churn',
    palette={'Retained': C_RETAIN, 'Churned': C_CHURN},
    plot_kws={'alpha': 0.25, 'edgecolor': 'none', 's': 15},
    diag_kind='kde',
    diag_kws={'alpha': 0.6}
)
pair_grid.fig.suptitle('Pair Plot — Top Predictive Features by Churn Status',
                        fontsize=14, fontweight='bold', y=1.01)
plt.show()
print('Day minutes clearly separates churned/retained along the x-axis in most panels.')
print('Customers with many vmail messages rarely churn (bottom-right panels).')

```



Day minutes clearly separates churned/retained along the x-axis in most panels.
 Customers with many vmail messages rarely churn (bottom-right panels).

2.6 Correlation Heatmap


Annotated heatmap showing pairwise correlations between numeric features — highlights collinearity between charge and minute columns.

```
In [13]: # VISUAL 6: Correlation heatmap
df_corr = df.copy()
df_corr['churn_int'] = df_corr['churn'].astype(int)
df_corr['intl_plan'] = (df_corr['international plan'] == 'yes').astype(int)
df_corr['vm_plan'] = (df_corr['voice mail plan'] == 'yes').astype(int)
drop_for_corr = ['state', 'phone number', 'area code', 'churn',
                 'international plan', 'voice mail plan']
corr_df = df_corr.drop(columns=drop_for_corr).corr()

# Shorten column labels
short = {
    'account length': 'Acct Len', 'number vmail messages': 'Vmail Msgs',
    'total day minutes': 'Day Min', 'total day calls': 'Day Calls',
    'total day charge': 'Day Charge', 'total eve minutes': 'Eve Min',
    'total eve calls': 'Eve Calls', 'total eve charge': 'Eve Charge',
    'total night minutes': 'Night Min', 'total night calls': 'Night Calls',
    'total night charge': 'Night Charge', 'total intl minutes': 'Intl Min',
    'total intl calls': 'Intl Calls', 'total intl charge': 'Intl Charge',
    'customer service calls': 'Svc Calls', 'churn_int': 'Churn',
    'intl_plan': 'Intl Plan', 'vm_plan': 'VM Plan'
}
corr_df = corr_df.rename(index=short, columns=short)

plt.figure(figsize=(16, 13))
mask = np.triu(np.ones_like(corr_df, dtype=bool))
sns.heatmap(
    corr_df, annot=True, fmt='.2f', cmap='RdBu_r',
    mask=mask, center=0, linewidths=0.5, vmin=-1, vmax=1,
    annot_kws={'size': 8},
    cbar_kws={'label': 'Pearson r', 'shrink': 0.8}
)
plt.title('Feature Correlation Heatmap', fontsize=14, fontweight='bold')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()

print('Charge columns have r=1.00 with their minute counterparts – perfect collinearity')
print('These will be dropped in Data Preparation to avoid redundancy.')
```



Charge columns have $r=1.00$ with their minute counterparts – perfect collinearity. These will be dropped in Data Preparation to avoid redundancy.

2.7 Data Understanding Summary

Finding	Implication
14.49% churn rate	Imbalanced — use recall & F1, not accuracy
International plan: 42.4% vs 11.5% churn	Strongest single predictor
4+ customer service calls: 45%+ churn	Non-linear threshold effect
Churned customers: 207 vs 175 day minutes	Higher usage → higher bills → sensitivity
Voice mail plan: 8.7% vs 16.7% churn	Protective factor / retention lever
Charge = $1.00 \times$ minutes	Drop charge columns (collinearity)
Zero missing values	No imputation needed

Phase 3: Data Preparation

```
In [14]: # Feature selection, encoding, and splitting

# Drop identifiers and collinear charge columns
drop_cols = [
    'phone number',      # unique identifier
    'state',              # 51 categories – high cardinality
    'area code',          # only 3 unique values, no signal
    'total day charge',    # r=1.00 with total day minutes
    'total eve charge',    # r=1.00 with total eve minutes
    'total night charge', # r=1.00 with total night minutes
    'total intl charge',  # r=1.00 with total intl minutes
]
df_model = df.drop(columns=drop_cols).copy()

# Encode binary categoricals
df_model['international plan'] = (df_model['international plan'] == 'yes').astype(int)
df_model['voice mail plan'] = (df_model['voice mail plan'] == 'yes').astype(int)
df_model['churn'] = df_model['churn'].astype(int)

# Stratified train / test split – 80/20
X = df_model.drop(columns='churn')
y = df_model['churn']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=RANDOM_STATE, stratify=y
)

# StandardScaler – fit on TRAIN only, transform both (prevents Leakage)
scaler = StandardScaler()
X_train_sc = scaler.fit_transform(X_train)
X_test_sc = scaler.transform(X_test)

print(f'Features retained: {X.shape[1]}')
print(f'Training set: {X_train.shape[0]:,} rows | churn rate: {y_train.mean()*100:.2f}')
print(f'Test set: {X_test.shape[0]:,} rows | churn rate: {y_test.mean()*100:.2f}')
print('\nScaler fitted on training data only – no data leakage.')
```

```
Features retained: 13
Training set: 2,666 rows | churn rate: 14.48%
Test set: 667 rows | churn rate: 14.54%
```

Scaler fitted on training data only – no data leakage.

3.1 Export Cleaned Dataset to CSV

After cleaning and encoding, `df_model` is saved as a CSV file so it can be inspected directly within the notebook and reused without re-running the full pipeline.

```
In [15]: # Export cleaned dataframe to CSV
output_path = 'churn_telecoms_cleaned.csv'
df_model.to_csv(output_path, index=False)

print(f'✅ Cleaned dataset saved to: {output_path}')
print(f'Shape: {df_model.shape[0]:,} rows × {df_model.shape[1]} columns')
print(f'Missing values: {df_model.isnull().sum().sum()}')
print()

# Preview the saved CSV by reading it back
df_preview = pd.read_csv(output_path)
```

```
print('--- Preview of cleaned dataset (first 5 rows) ---')
df_preview.head()
```

✔ Cleaned dataset saved to: churn_telecoms_cleaned.csv
Shape: 3,333 rows × 14 columns
Missing values: 0

--- Preview of cleaned dataset (first 5 rows) ---

Out[15]:

	account length	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total eve minutes	total eve calls	total night minutes	total night calls	total intl minutes	total intl calls
0	128	0	1	25	265.1	110	197.4	99	244.7	91	10.0	
1	107	0	1	26	161.6	123	195.5	103	254.4	103	13.7	
2	137	0	0	0	243.4	114	121.2	110	162.6	104	12.2	
3	84	1	0	0	299.4	71	61.9	88	196.9	89	6.6	
4	75	1	0	0	166.7	113	148.3	122	186.9	121	10.1	

Phase 4: Modeling

Strategy

Model	Type	Key Change from Previous
M1	Logistic Regression (Baseline)	Default — performance floor
M2	Logistic Regression (Tuned)	<code>class_weight='balanced'</code> + GridSearch on C
M3	Decision Tree (Baseline)	Non-parametric — tests non-linear patterns
M4	Decision Tree (Tuned)	GridSearch on depth/leaf — reduces overfitting
M5	Random Forest (Final)	Ensemble — reduces variance, best generalisation

```
In [16]: # — Evaluation helper —————
def evaluate_model(model, Xtr, ytr, Xte, yte, name):
    """Fit model, return metrics dict, print report, show confusion matrix."""
    model.fit(Xtr, ytr)
    tr_pred = model.predict(Xtr)
    te_pred = model.predict(Xte)

    tr_rec = recall_score(ytr, tr_pred)
    te_rec = recall_score(yte, te_pred)
    tr_f1 = f1_score(ytr, tr_pred)
    te_f1 = f1_score(yte, te_pred)
    te_pre = precision_score(yte, te_pred)
    te_auc = roc_auc_score(yte, model.predict_proba(Xte)[:, 1])

    print(f'\n{"="*60}')
    print(f' {name}')
    print(f'{"="*60}')
    print(f' Recall    - Train: {tr_rec:.3f} | Test: {te_rec:.3f}')
```



```

print(f' F1-Score - Train: {tr_f1:.3f} | Test: {te_f1:.3f}')
print(f' ROC-AUC (Test): {te_auc:.3f}')
print(f'\n Classification Report (Test):')
print(classification_report(yte, te_pred, target_names=['No Churn', 'Churn']))

fig, ax = plt.subplots(figsize=(5, 4))
ConfusionMatrixDisplay.from_predictions(
    yte, te_pred, display_labels=['No Churn', 'Churn'],
    cmap='Blues', ax=ax, colorbar=False
)
ax.set_title(f'Confusion Matrix\n{name}', fontweight='bold', fontsize=11)
plt.tight_layout(); plt.show()

return dict(Model=name,
            Train_Recall=round(tr_rec,3), Test_Recall=round(te_rec,3),
            Train_F1=round(tr_f1,3), Test_F1=round(te_f1,3),
            Test_Precision=round(te_pre,3), Test_AUC=round(te_auc,3))

all_results = []
all_models = [] # store (model, Xte, name) for ROC/PR curves

```

Model 1 — Logistic Regression (Baseline)

Rationale: Establishes performance floor. Logistic regression is fast, interpretable, and requires feature scaling. Default parameters — no class imbalance correction — so we can see the raw imbalance impact.

```

In [17]: lr_base = LogisticRegression(max_iter=1000, random_state=RANDOM_STATE)
res = evaluate_model(lr_base, X_train_sc, y_train,
                    X_test_sc, y_test, 'M1: Logistic Regression (Baseline)')
all_results.append(res)
all_models.append((lr_base, X_test_sc, 'M1: LR Baseline'))

```

```


=====
M1: Logistic Regression (Baseline)
=====
Recall    - Train: 0.207 | Test: 0.247
F1-Score  - Train: 0.302 | Test: 0.338
ROC-AUC   (Test):  0.817

Classification Report (Test):
      precision    recall  f1-score   support

No Churn      0.88      0.96      0.92       570
Churn         0.53      0.25      0.34        97

accuracy              0.86       667
macro avg      0.71      0.61      0.63       667
weighted avg   0.83      0.86      0.84       667

```



Interpretation: Only **24.7% recall** — the model ignores the minority class due to class imbalance. ROC-AUC of 0.817 shows the model *can* rank customers, but the default threshold of 0.5 is badly calibrated. Next step: correct for imbalance.

Model 2 — Logistic Regression (Tuned: class_weight + GridSearchCV)

Rationale: Add `class_weight='balanced'` to penalise misclassifying churners more heavily. Use GridSearchCV on regularisation strength `C` (5-fold, scoring=recall) to find the best regularisation level.

```
In [18]: gs_lr = GridSearchCV(
    LogisticRegression(max_iter=1000, random_state=RANDOM_STATE, class_weight='balanced'
    {'C': [0.01, 0.1, 1, 10, 100]}, cv=5, scoring='recall', n_jobs=-1
)
gs_lr.fit(X_train_sc, y_train)
print(f'Best C: {gs_lr.best_params_["C"]} | CV Recall: {gs_lr.best_score_:.3f}')

lr_tuned = gs_lr.best_estimator_
res = evaluate_model(lr_tuned, X_train_sc, y_train,
                    X_test_sc, y_test, 'M2: Logistic Regression (Balanced, Tuned)')
all_results.append(res)
all_models.append((lr_tuned, X_test_sc, 'M2: LR Tuned'))
```

Best C: 1 | CV Recall: 0.752


```
=====
M2: Logistic Regression (Balanced, Tuned)
```

```
=====
Recall    - Train: 0.764 | Test: 0.742
F1-Score  - Train: 0.494 | Test: 0.472
ROC-AUC   (Test):  0.815
```

```
Classification Report (Test):
              precision    recall  f1-score   support

   No Churn         0.95         0.76         0.84         570
     Churn          0.35         0.74         0.47          97

 accuracy                   0.76         667
 macro avg                 0.65         0.75         0.66         667
 weighted avg              0.86         0.76         0.79         667
```



Interpretation: Recall jumps to **74.2%** — a 50-point improvement. Precision drops (35%) because the model now flags more non-churners too. F1 of 0.474 reflects the trade-off. ROC-AUC unchanged — the model's ranking ability was always there, the threshold was the problem. Next: try non-linear models that may maintain recall while recovering precision.

Model 3 — Decision Tree (Baseline)

Rationale: Logistic regression draws linear boundaries. Churn drivers are non-linear (e.g., churn spikes sharply after 4+ service calls — not a linear effect). Decision Trees can capture these threshold effects and interactions without feature scaling.

```
In [19]: dt_base = DecisionTreeClassifier(random_state=RANDOM_STATE)
         res = evaluate_model(dt_base, X_train, y_train,
```


```
X_test, y_test, 'M3: Decision Tree (Baseline)')
all_results.append(res)
all_models.append((dt_base, X_test, 'M3: DT Baseline'))
```

```
=====
M3: Decision Tree (Baseline)
=====
```

```
Recall    - Train: 1.000 | Test: 0.649
F1-Score  - Train: 1.000 | Test: 0.663
ROC-AUC   (Test):  0.798
```

```
Classification Report (Test):
```

	precision	recall	f1-score	support
No Churn	0.94	0.95	0.94	570
Churn	0.68	0.65	0.66	97
accuracy			0.90	667
macro avg	0.81	0.80	0.80	667
weighted avg	0.90	0.90	0.90	667



Interpretation: Train recall = 1.00, Test recall = 0.649 — textbook overfitting. The tree memorises the training data. Despite this, F1 (0.663) already beats Model 2 (0.474), confirming that the non-linear approach is better suited to this data. Fix: constrain tree depth.

Model 4 — Decision Tree (Tuned)

Rationale: Model 3 overfits heavily (train-test gap = 0.351). Tune `max_depth`, `min_samples_leaf`, and `class_weight` via GridSearchCV (5-fold, scoring=recall) to constrain

the tree and improve generalisation.

```
In [20]: gs_dt = GridSearchCV(
    DecisionTreeClassifier(random_state=RANDOM_STATE),
    {'max_depth': [3, 5, 7, 10, None],
     'min_samples_leaf': [1, 5, 10],
     'class_weight': ['balanced', None]},
    cv=5, scoring='recall', n_jobs=-1
)
gs_dt.fit(X_train, y_train)
print(f'Best params: {gs_dt.best_params_}')
print(f'CV Recall: {gs_dt.best_score_:.3f}')

dt_tuned = gs_dt.best_estimator_
res = evaluate_model(dt_tuned, X_train, y_train,
                     X_test, y_test, 'M4: Decision Tree (Tuned)')
all_results.append(res)
all_models.append((dt_tuned, X_test, 'M4: DT Tuned'))
```


```
Best params: {'class_weight': 'balanced', 'max_depth': 10, 'min_samples_leaf': 5}
CV Recall: 0.811
```

```
=====
M4: Decision Tree (Tuned)
=====
Recall    - Train: 0.938 | Test: 0.742
F1-Score  - Train: 0.852 | Test: 0.673
ROC-AUC   (Test): 0.833


Classification Report (Test):
      precision    recall  f1-score   support

No Churn      0.95      0.92      0.94       570
Churn         0.62      0.74      0.67        97

accuracy              0.90       667
macro avg      0.78      0.83      0.81       667
weighted avg   0.91      0.90      0.90       667
```



```
In [21]: # Visualise tuned decision tree (top 3 levels)
fig, ax = plt.subplots(figsize=(22, 8))
plot_tree(dt_tuned, feature_names=X_train.columns.tolist(),
          class_names=['No Churn', 'Churn'], filled=True,
          max_depth=3, fontsize=9, ax=ax, impurity=False)
plt.title('Tuned Decision Tree — Top 3 Levels', fontsize=14, fontweight='bold')
plt.tight_layout(); plt.show()
```



Interpretation: Overfitting gap shrinks from 0.351 → 0.196. Test recall rises to **74.2%** matching M2, but F1 is better (0.673 vs 0.474) — same recall with far less collateral. The tree first splits on **customer service calls** and **total day minutes**, confirming the EDA insights.

Model 5 — Random Forest (Final Model)

Rationale: Model 4 still overfits (train 0.938 vs test 0.742). A Random Forest:

- Trains 200 trees on bootstrap samples (reduces variance via bagging)
- Each split considers only a random subset of features (de-correlates trees)
- Averages predictions — individual tree errors cancel out

Expected: best generalisation, highest ROC-AUC, most stable feature importance estimates.

```
In [22]: gs_rf = GridSearchCV(
    RandomForestClassifier(random_state=RANDOM_STATE, n_jobs=-1),
    {'n_estimators': [100, 200],
     'max_depth':    [5, 10, None],
     'class_weight': ['balanced']},
    cv=5, scoring='recall', n_jobs=-1
  )
gs_rf.fit(X_train, y_train)
print(f'Best params: {gs_rf.best_params_}')
print(f'CV Recall:    {gs_rf.best_score_:.3f}')


rf_final = gs_rf.best_estimator_
res = evaluate_model(rf_final, X_train, y_train,
                    X_test, y_test, 'M5: Random Forest (Final)')
all_results.append(res)
all_models.append((rf_final, X_test, 'M5: RF Final ★'))
```

```
Best params: {'class_weight': 'balanced', 'max_depth': 5, 'n_estimators': 200}
CV Recall:    0.819
```

```
=====
M5: Random Forest (Final)
=====
Recall    - Train: 0.852 | Test: 0.722
F1-Score  - Train: 0.759 | Test: 0.622
ROC-AUC   (Test): 0.882

Classification Report (Test):
```

	precision	recall	f1-score	support
No Churn	0.95	0.90	0.92	570
Churn	0.55	0.72	0.62	97
accuracy			0.87	667
macro avg	0.75	0.81	0.77	667
weighted avg	0.89	0.87	0.88	667



Interpretation: The Random Forest achieves the **smallest train-test gap (0.130)** and highest **ROC-AUC (0.882)** of all five models. Test recall of 0.722 and F1 of 0.622 confirm strong generalisation. This is the selected final model.

Phase 5: Evaluation

5.1 Model Comparison — Dumbbell Chart (Train vs Test Recall)

A dumbbell chart elegantly shows the train-test recall gap for each model — the distance between the two dots reveals how much each model overfits.

```
In [23]: # VISUAL 7: Dumbbell chart – train vs test recall gap
results_df = pd.DataFrame(all_results)
short_names = ['M1: LR\nBaseline', 'M2: LR\nTuned', 'M3: DT\nBaseline',
               'M4: DT\nTuned', 'M5: RF\nFinal ★']

fig, ax = plt.subplots(figsize=(13, 6))
y_pos = range(len(short_names))

for i, (_, row) in enumerate(results_df.iterrows()):
    tr = row['Train_Recall']
    te = row['Test_Recall']
    color = C_PURPLE if i == 4 else '#555555'
    lw = 3 if i == 4 else 1.8
    ax.plot([tr, te], [i, i], color=color, linewidth=lw, zorder=2)
```




```

ax.plot(tr, i, 'o', color=C_BLUE, markersize=13, zorder=3,
        label='Train Recall' if i == 0 else '')
ax.plot(te, i, 's', color=C_CHURN, markersize=13, zorder=3,
        label='Test Recall' if i == 0 else '')
ax.text(tr - 0.013, i, f'{tr:.3f}', va='center', ha='right',
        fontsize=9, color=C_BLUE, fontweight='bold')
ax.text(te + 0.013, i, f'{te:.3f}', va='center', ha='left',
        fontsize=9, color=C_CHURN, fontweight='bold')
gap = round(tr - te, 3)
ax.text((tr+te)/2, i + 0.22, f'gap={gap:.3f}',
        ha='center', fontsize=8, color='#888888')

ax.set_yticks(y_pos)
ax.set_yticklabels(short_names, fontsize=11)
ax.set_xlabel('Recall Score', fontsize=12)
ax.set_xlim(0.05, 1.10)
ax.axvline(x=0.70, color='navy', linestyle='--', linewidth=1.2,
           label='Business target recall ≥ 0.70')
ax.set_title('Train vs Test Recall — Overfitting Gap by Model\n'
            '(Highlighted row = selected final model)',
            fontsize=13, fontweight='bold')
ax.legend(fontsize=10, loc='lower right')

# Highlight final model row
ax.axhspan(3.5, 4.5, alpha=0.08, color=C_PURPLE)
plt.tight_layout()
plt.show()

```



5.2 Radar / Spider Chart — All Metrics Across All Models

A radar chart compares all five models across four metrics simultaneously, making it easy to spot well-rounded models vs those that excel on one metric at the expense of others.

```

In [24]: # — VISUAL 8: Radar chart — all models, all metrics —————
categories = ['Test\nRecall', 'Test\nF1', 'Test\nAUC', 'Test\nPrecision']
metric_keys = ['Test_Recall', 'Test_F1', 'Test_AUC', 'Test_Precision']
N = len(categories)
angles = [n / float(N) * 2 * np.pi for n in range(N)]
angles += angles[:1]

```


```
radar_colors = [C_BLUE, C_ORANGE, C_CHURN, C_RETAIN, C_PURPLE]
lw_vals = [1.5, 1.5, 1.5, 1.5, 2.8]
alpha_vals = [0.05, 0.05, 0.05, 0.05, 0.15]

fig, ax = plt.subplots(figsize=(9, 9), subplot_kw=dict(polar=True))

for i, (_, row) in enumerate(results_df.iterrows()):
    values = [row[k] for k in metric_keys]
    values += values[:1]
    ax.plot(angles, values, color=radar_colors[i],
            linewidth=lw_vals[i], linestyle='solid',
            label=short_names[i].replace('\n', ' '))
    ax.fill(angles, values, color=radar_colors[i], alpha=alpha_vals[i])

ax.set_xticks(angles[:-1])
ax.set_xticklabels(categories, fontsize=12, fontweight='bold')
ax.set_ylim(0, 1)
ax.set_yticks([0.2, 0.4, 0.6, 0.8, 1.0])
ax.set_yticklabels(['0.2', '0.4', '0.6', '0.8', '1.0'], fontsize=8, color='grey')
ax.set_title('Model Comparison – Radar Chart\n(All metrics on test set)',
            fontsize=14, fontweight='bold', pad=20)
ax.legend(loc='upper right', bbox_to_anchor=(1.35, 1.15), fontsize=10)
ax.grid(color='grey', linestyle='--', linewidth=0.5, alpha=0.5)

plt.tight_layout()
plt.show()
```



5.3 ROC Curves — All Models

ROC curves show the full trade-off between true positive rate (recall) and false positive rate at every possible probability threshold. Higher and further left = better.


```
In [25]: # VISUAL 9: ROC Curves
roc_colors = [C_BLUE, C_ORANGE, C_CHURN, C_RETAIN, C_PURPLE]
line_styles = ['--', '-.', ':', (0,(3,1,1,1)), '-']
lw_roc = [1.5, 1.5, 1.5, 1.5, 2.5]

fig, ax = plt.subplots(figsize=(9, 8))

for (model, Xte, label), color, ls, lw in zip(all_models, roc_colors, line_styles, lw_roc):
    proba = model.predict_proba(Xte)[:, 1]
    fpr, tpr, _ = roc_curve(y_test, proba)
    auc = roc_auc_score(y_test, proba)
    ax.plot(fpr, tpr, color=color, linestyle=ls, linewidth=lw,
            label=f'{label} (AUC = {auc:.3f})')

ax.plot([0,1],[0,1], 'k--', linewidth=1, label='Random guess (AUC = 0.500)')
ax.fill_between([0,1],[0,1], alpha=0.05, color='grey')
ax.set_xlabel('False Positive Rate (1 - Specificity)', fontsize=12)
ax.set_ylabel('True Positive Rate (Recall / Sensitivity)', fontsize=12)
```

```
ax.set_title('ROC Curves — All Models\n(M5 Random Forest achieves highest AUC)',
            fontsize=13, fontweight='bold')
ax.legend(fontsize=10, loc='lower right')
ax.set_xlim([0, 1]); ax.set_ylim([0, 1.02])
plt.tight_layout(); plt.show()
```



5.4 Precision-Recall Curve — Final Model

For imbalanced datasets, the Precision-Recall (PR) curve is more informative than the ROC curve. It shows every recall level's achievable precision — the business can choose where on this curve to operate depending on team capacity.

```
In [26]: # VISUAL 10: Precision-Recall Curve — Final Model
y_proba_rf = rf_final.predict_proba(X_test)[: , 1]
precision_vals, recall_vals, thresholds = precision_recall_curve(y_test, y_proba_rf)
avg_precision = average_precision_score(y_test, y_proba_rf)

fig, ax = plt.subplots(figsize=(10, 7))

# PR curve coloured by threshold
sc = ax.scatter(recall_vals[:-1], precision_vals[:-1],
               c=thresholds, cmap='RdYlGn_r', s=8, zorder=3)
ax.plot(recall_vals, precision_vals, color='grey', linewidth=1.2,
```

```

alpha=0.5, zorder=2)
plt.colorbar(sc, ax=ax, label='Decision Threshold', shrink=0.8)


# Mark default threshold (0.5)
idx_50 = np.argmin(np.abs(thresholds - 0.5))
ax.plot(recall_vals[idx_50], precision_vals[idx_50],
        'b*', markersize=18, label=f'Threshold=0.50 (Recall={recall_vals[idx_50]:.2f},
        zorder=5)

# No-skill baseline
no_skill = y_test.mean()
ax.axhline(y=no_skill, color='navy', linestyle='--', linewidth=1.2,
           label=f'No-skill baseline (precision = {no_skill:.2f})')

ax.set_xlabel('Recall (True Positive Rate)', fontsize=12)
ax.set_ylabel('Precision', fontsize=12)
ax.set_title(f'Precision-Recall Curve — Random Forest Final Model\n'
            f'Average Precision = {avg_precision:.3f}',
            fontsize=13, fontweight='bold')
ax.legend(fontsize=10)
ax.set_xlim([0, 1]); ax.set_ylim([0, 1.05])
plt.tight_layout(); plt.show()

print('The colour gradient shows how the threshold changes along the curve.')
print('Moving right (more recall) requires lowering the threshold → accepting lower pre')
print('The retention team can move the operating point based on their outreach capacity')

```



The colour gradient shows how the threshold changes along the curve.
Moving right (more recall) requires lowering the threshold → accepting lower precision.
The retention team can move the operating point based on their outreach capacity.

5.5 Feature Importance — Lollipop Chart

A lollipop chart conveys the same ranked importance information as a bar chart but with much less visual clutter — the thin stem draws the eye to the dot without dominating the figure.


```
In [27]: # VISUAL 11: Lollipop chart – feature importance
importances = pd.Series(rf_final.feature_importances_, index=X_train.columns)
importances = importances.sort_values(ascending=True)
top_q = importances.quantile(0.70)

fig, ax = plt.subplots(figsize=(11, 8))

for i, (feat, val) in enumerate(importances.items()):
    color = C_CHURN if val >= top_q else C_BLUE
    ax.hlines(i, 0, val, linewidth=2, color=color, alpha=0.6)
    ax.plot(val, i, 'o', markersize=12, color=color, zorder=5)
    ax.text(val + 0.003, i, f'{val:.4f}',
            va='center', fontsize=9, color=color, fontweight='bold')

ax.axvline(x=top_q, color='navy', linestyle='--', linewidth=1.5,
           label=f'70th percentile ({top_q:.4f})')
ax.set_yticks(range(len(importances)))
ax.set_yticklabels(importances.index, fontsize=10)
ax.set_xlabel('Mean Decrease in Impurity (Feature Importance)', fontsize=11)
ax.set_title('Feature Importances – Random Forest Final Model\n'
            '(Red = top 30% most important features)',
            fontsize=13, fontweight='bold')
ax.legend(fontsize=10)
ax.set_xlim(0, importances.max() + 0.04)
plt.tight_layout(); plt.show()

print('Top 5 features:')
for feat, val in importances.sort_values(ascending=False).head(5).items():
    print(f' {feat:<30s} {val:.4f}')
```



Top 5 features:

total day minutes	0.3033
customer service calls	0.2289
international plan	0.1539
total eve minutes	0.0615
number vmail messages	0.0502

5.6 Feature Importance — Treemap

A treemap presents feature importances as proportional rectangles — making the relative magnitude of each feature immediately visible at a glance.

```
In [28]: # VISUAL 12: Treemap — feature importance
import squarify

imp_sorted = importances.sort_values(ascending=False)

# Colour by importance tier
def tier_color(v, max_v):
    ratio = v / max_v
    if ratio > 0.50: return '#c0392b'
    if ratio > 0.30: return '#e67e22'
    if ratio > 0.15: return '#3498db'
    return '#95a5a6'

colors_tm = [tier_color(v, imp_sorted.max()) for v in imp_sorted.values]
labels_tm = [f'{name}\n{val:.3f}' for name, val in imp_sorted.items()]


fig, ax = plt.subplots(figsize=(14, 8))
squarify.plot(
    sizes=imp_sorted.values,
```

```

label=labels_tm,
color=colors_tm,
alpha=0.88,
text_kwargs={'fontsize': 10, 'fontweight': 'bold', 'color': 'white'},
ax=ax,
pad=True
)
ax.axis('off')
ax.set_title('Feature Importance Treemap — Random Forest Final Model\n'
            'Area proportional to importance | Red = highest, grey = lowest',
            fontsize=13, fontweight='bold')

# Legend
legend_patches = [
    mpatches.Patch(color='#c0392b', label='>50% of max importance'),
    mpatches.Patch(color='#e67e22', label='30-50%'),
    mpatches.Patch(color='#3498db', label='15-30%'),
    mpatches.Patch(color='#95a5a6', label='<15%'),
]
ax.legend(handles=legend_patches, loc='lower right', fontsize=9, frameon=True)
plt.tight_layout(); plt.show()

```



5.7 Final Model — Holdout Test Set Evaluation

```

In [29]: # Final model metrics on holdout test set
y_pred_final = rf_final.predict(X_test)
y_proba_final = rf_final.predict_proba(X_test)[:, 1]
tn, fp, fn, tp = confusion_matrix(y_test, y_pred_final).ravel()

print('FINAL MODEL — HOLDOUT TEST SET RESULTS')
print('='*55)
print(f'Model:          Random Forest (n_estimators=200, max_depth=5, class_weight=balan
print(f'Test size:      {len(y_test):,} customers ({y_test.sum():,} churned | {len(y_test)
print()
print(f' True Positives (churners caught):          {tp:>4}')
print(f' False Negatives (churners missed):         {fn:>4}')

```



```

print(f' False Positives (non-churners flagged):      {fp:>4}')
print(f' True Negatives (correctly cleared):          {tn:>4}')
print()
print(f' Recall      : {recall_score(y_test,y_pred_final):.3f}    ({tp}/{y_test.sum()} ch
print(f' Precision  : {precision_score(y_test,y_pred_final):.3f}')
print(f' F1-Score   : {f1_score(y_test,y_pred_final):.3f}')
print(f' ROC-AUC    : {roc_auc_score(y_test,y_proba_final):.3f}')
print()
print('Business translation:')
print(f' Of {y_test.sum()} customers who would churn, the model correctly flags {tp} (
print(f' The retention team handles {fp} false alarms alongside {tp} real churners.')
print(f' Roughly {precision_score(y_test,y_pred_final)*100:.0f}% of all flagged custom

```

FINAL MODEL – HOLDOUT TEST SET RESULTS

```

=====
Model:      Random Forest (n_estimators=200, max_depth=5, class_weight=balanced)
Test size:   667 customers (97 churned | 570 retained)

```

```

True Positives (churners caught):      70
False Negatives (churners missed):      27
False Positives (non-churners flagged):  58
True Negatives (correctly cleared):     512

```

```

Recall      : 0.722    (70/97 churners identified)
Precision   : 0.547
F1-Score    : 0.622
ROC-AUC     : 0.882

```

Business translation:

Of 97 customers who would churn, the model correctly flags 70 (72%).
 The retention team handles 58 false alarms alongside 70 real churners.
 Roughly 55% of all flagged customers are genuine churn risks.

5.8 Final Model Selection Justification

Criterion	M4: Decision Tree	M5: Random Forest	Verdict
Test Recall	0.742	0.722	Within noise — tie
Test F1	0.673	0.622	DT slightly higher
Test AUC	0.833	0.882	RF significantly higher
Train-test gap	0.196	0.130	RF generalises better

M5 (Random Forest) selected because:

1. **ROC-AUC 0.882** — the retention team scores customers by *probability*, not binary predictions. AUC directly measures ranking quality at every threshold.
2. **Smallest overfitting gap (0.130)** — more reliable on truly unseen production data.
3. **Ensemble feature importances** — stable estimates averaged across 200 trees rather than a single tree's arbitrary splits.

Phase 6: Conclusions, Recommendations & Deployment

Key Feature Insights

Rank	Feature	Business Implication	Recommended Action
1	Total day minutes	Heavy users pay more → price-sensitive → higher churn	Offer loyalty rate caps for customers above 75th percentile (>216 min/day)
2	Customer service calls	4+ calls signals deep dissatisfaction	Auto-flag after 3rd call; route to senior retention agent
3	International plan	42.4% vs 11.5% churn — plan may be poor value	Audit pricing and coverage; conduct exit interviews
4	Total eve minutes	Secondary usage signal	Monitor alongside day minutes
5	Total intl minutes	High intl usage + intl plan = double churn risk	Offer discounted intl bundles to heavy intl users

Protective factor: Voice mail plan holders churn at 8.7% vs 16.7% — consider offering it as a free retention incentive.

Model Limitations

1. **No demographic or contract data** — age, income, and contract type are known churn drivers but absent.
2. **No timestamps** — cannot model churn as a time-to-event problem or detect seasonality.
3. **Static model** — must be retrained as pricing and competitive landscape evolve.
4. **27.8% missed churners** — the model still misses roughly 1 in 4 actual churners.
5. **45% false positive rate among flagged** — outreach cost for non-churners must be budgeted.

Deployment Recommendations

- **Weekly scoring** — score all active accounts weekly; export priority-ranked list to CRM.
- **Adjustable threshold** — lower the probability threshold to catch more churners; raise it to reduce false alarms based on team capacity.
- **Quarterly retraining** — refresh the model on new data every quarter.
- **A/B test** — hold out a random control group of high-risk customers to measure true retention lift from model-driven outreach.